

```
In [1]: pip install pybaseball
```

```
Requirement already satisfied: pybaseball in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (2.2.7)
Requirement already satisfied: numpy>=1.13.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (1.26.1)
Requirement already satisfied: pandas>=1.0.3 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (2.1.2)
Requirement already satisfied: beautifulsoup4>=4.4.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (4.12.2)
Requirement already satisfied: requests>=2.18.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (2.31.0)
Requirement already satisfied: lxml>=4.2.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (4.9.3)
Requirement already satisfied: pyarrow>=1.0.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (14.0.0)
Requirement already satisfied: pygithub>=1.51 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (2.1.1)
Requirement already satisfied: scipy>=1.4.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (1.11.3)
Requirement already satisfied: matplotlib>=2.0.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (3.8.1)
Requirement already satisfied: tqdm>=4.50.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (4.66.1)
Requirement already satisfied: attrs>=20.3.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pybaseball) (23.1.0)
Requirement already satisfied: soupsieve>1.2 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from beautifulsoup4>=4.4.0->pybaseball) (2.5)
Requirement already satisfied: contourpy>=1.0.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (1.1.1)
Requirement already satisfied: cycler>=0.10 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (4.43.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (1.4.5)
```

Requirement already satisfied: packaging>=20.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (23.2)

Requirement already satisfied: pillow>=8 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (10.1.0)

Requirement already satisfied: pyparsing>=2.3.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (3.1.1)

Requirement already satisfied: python-dateutil>=2.7 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (2.8.2)

Requirement already satisfied: importlib-resources>=3.2.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from matplotlib>=2.0.0->pybaseball) (6.1.0)

Requirement already satisfied: pytz>=2020.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pandas>=1.0.3->pybaseball) (2023.3.post1)

Requirement already satisfied: tzdata>=2022.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pandas>=1.0.3->pybaseball) (2023.3)

Requirement already satisfied: pynacl>=1.4.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pygithub>=1.51->pybaseball) (1.5.0)

Requirement already satisfied: pyjwt>=2.4.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pyjwt[crypto]>=2.4.0->pygithub>=1.51->pybaseball) (2.8.0)

Requirement already satisfied: typing-extensions>=4.0.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pygithub>=1.51->pybaseball) (4.8.0)

Requirement already satisfied: urllib3>=1.26.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pygithub>=1.51->pybaseball) (2.0.7)

Requirement already satisfied: Deprecated in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pygithub>=1.51->pybaseball) (1.2.14)

Requirement already satisfied: charset-normalizer<4,>=2 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from requests>=2.18.1->pybaseball) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from requests>=2.18.1->pybaseball) (3.4)

Requirement already satisfied: certifi>=2017.4.17 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from requests>=2.18.1->pybaseball) (2023.7.22)

Requirement already satisfied: zipp>=3.1.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from importlib-resources>=3.2.0->matplotlib>=2.0.0->pybaseball) (3.17.0)

Requirement already satisfied: cryptography>=3.4.0 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from pyjwt[crypto]>=2.4.0->pygithub>=1.51->pybaseball) (41.0.5)

Requirement already satisfied: cffi>=1.4.1 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from

pynacl>=1.4.0->pygithub>=1.51->pybaseball) (1.16.0)
Requirement already satisfied: six>=1.5 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from python-dateutil>=2.7->matplotlib>=2.0.0->pybaseball) (1.16.0)
Requirement already satisfied: wrapt<2,>=1.10 in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from Deprecated->pygithub>=1.51->pybaseball) (1.14.1)
Requirement already satisfied: pycparser in /Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages (from cffi>=1.4.1->pynacl>=1.4.0->pygithub>=1.51->pybaseball) (2.21)
Note: you may need to restart the kernel to use updated packages.

In [2]: *#Import initial libraries*

```
import pybaseball
import numpy as np
import pandas as pd
import matplotlib.pyplot as plot
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
```

/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages/urllib3/__init__.py:34: NotOpenSSLWarning: urllib3 v2.0 only supports OpenSSL 1.1.1+, currently the 'ssl' module is compiled with 'LibreSSL 2.8.3'. See: <https://github.com/urllib3/urllib3/issues/3020>
warnings.warn(

In [3]: *#We use pybaseball data and split it in regular and post season dataframes*
df_reg = pybaseball.statcast(start_dt = '2023-03-30', end_dt = '2023-10-02')
df_post = pybaseball.statcast(start_dt = '2023-10-03', end_dt = '2023-11-01')

This is a large query, it may take a moment to complete

```

/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/pybaseball/statcast.py:50: UserWarning:
That's a nice request you got there. It'd be a shame if something were to ha
ppen to it.
We strongly recommend that you enable caching before running this. It's as s
imple as `pybaseball.cache.enable()`.
Since the Statcast requests can take a *really* long time to run, if somethi
ng were to happen, like: a disconnect;
gremlins; computer repair by associates of Rudy Giuliani; electromagnetic in
terference from metal trash cans; etc.;
you could lose a lot of progress. Enabling caching will allow you to immedia
tely recover all the successful
subqueries if that happens.
    warnings.warn(_OVERSIZE_WARNING)
100%|████████████████████████████████████████████████████████████████████████████████| 187/187 [01:28<00:00, 2.11i
t/s]
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/pybaseball/statcast.py:85: FutureWarning: The behavi
or of DataFrame concatenation with empty or all-NA entries is deprecated. In
a future version, this will no longer exclude empty or all-NA columns when d
etermining the result dtypes. To retain the old behavior, exclude the releva
nt entries before the concat operation.
    final_data = pd.concat(dataframe_list, axis=0).convert_dtypes(convert_stri
ng=False)

```

This is a large query, it may take a moment to complete

```

100%|████████████████████████████████████████████████████████████████████████████████| 30/30 [00:04<00:00, 6.18i
t/s]
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/pybaseball/statcast.py:85: FutureWarning: The behavi
or of DataFrame concatenation with empty or all-NA entries is deprecated. In
a future version, this will no longer exclude empty or all-NA columns when d
etermining the result dtypes. To retain the old behavior, exclude the releva
nt entries before the concat operation.
    final_data = pd.concat(dataframe_list, axis=0).convert_dtypes(convert_stri
ng=False)

```

In [4]: *#We handpicked columns that we felt did not convey much information or were
#After this, we ended up with 32 columns that we will use in our analyses*

```

df_reg = df_reg.drop(['game_date', 'spin_dir', 'spin_rate_deprecated', 'brea
    'des', 'game_year', 'inning_topbot', 'hc_x', 'hc_y', 'tfs_c
    'umpire', 'sv_id', 'vx0', 'vy0', 'vz0', 'ax', 'ay', 'az', '
    'release_extension', 'release_pos_y', 'babip_value', 'woba
    'bat_score', 'fld_score', 'post_away_score', 'post_home_scc
    'spin_axis', 'release_speed', 'events', 'description', 'zor
    'pfx_x', 'pfx_z', 'plate_x', 'plate_z', 'hit_distance_sc',
    'estimated_woba_using_speedangle', 'launch_speed_angle', '
df_post = df_post.drop(['game_date', 'spin_dir', 'spin_rate_deprecated', 'br
    'des', 'game_year', 'inning_topbot', 'hc_x', 'hc_y', 'tfs_c
    'umpire', 'sv_id', 'vx0', 'vy0', 'vz0', 'ax', 'ay', 'az', '
    'release_extension', 'release_pos_y', 'babip_value', 'woba
    'bat_score', 'fld_score', 'post_away_score', 'post_home_scc
    'spin_axis', 'release_speed', 'events', 'description', 'zor

```

```
'pfx_x', 'pfx_z', 'plate_x', 'plate_z', 'hit_distance_sc',  
'estimated_woba_using_speedangle', 'launch_speed_angle', '
```

```
In [5]: df_reg.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 717945 entries, 2437 to 4497  
Data columns (total 34 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   release_pos_x                        717673 non-null  Float64  
1   release_pos_z                        717673 non-null  Float64  
2   player_name                          717945 non-null  object  
3   batter                              717945 non-null  Int64  
4   pitcher                              717945 non-null  Int64  
5   stand                               717945 non-null  object  
6   p_throws                             717945 non-null  object  
7   home_team                            717945 non-null  object  
8   away_team                            717945 non-null  object  
9   type                                 717945 non-null  object  
10  balls                                717945 non-null  Int64  
11  strikes                              717945 non-null  Int64  
12  on_3b                                69419 non-null   Int64  
13  on_2b                                137481 non-null  Int64  
14  on_1b                                221293 non-null  Int64  
15  outs_when_up                         717945 non-null  Int64  
16  inning                               717945 non-null  Int64  
17  game_pk                              717945 non-null  Int64  
18  pitcher.1                            717945 non-null  Int64  
19  fielder_2.1                          717945 non-null  Int64  
20  fielder_3                            717945 non-null  Int64  
21  fielder_4                            717945 non-null  Int64  
22  fielder_5                            717945 non-null  Int64  
23  fielder_6                            717945 non-null  Int64  
24  fielder_7                            717945 non-null  Int64  
25  fielder_8                            717945 non-null  Int64  
26  fielder_9                            717945 non-null  Int64  
27  at_bat_number                        717945 non-null  Int64  
28  pitch_number                         717945 non-null  Int64  
29  pitch_name                           717677 non-null  object  
30  home_score                           717945 non-null  Int64  
31  away_score                           717945 non-null  Int64  
32  if_fielding_alignment                715432 non-null  object  
33  of_fielding_alignment                715432 non-null  object  
dtypes: Float64(2), Int64(23), object(9)  
memory usage: 208.8+ MB
```

```
In [6]: #We use 14 pitchers that will also be playing in the postseason to filter out  
  
df_reg_filtered = df_reg[df_reg['pitcher'].isin([605400, 554430, 624133, 543133, 543134, 543135, 543136, 543137, 543138, 543139, 543140, 543141, 543142, 543143])]  
df_post_filtered = df_post[df_post['pitcher'].isin([605400, 554430, 624133, 543133, 543134, 543135, 543136, 543137, 543138, 543139, 543140, 543141, 543142, 543143])]
```

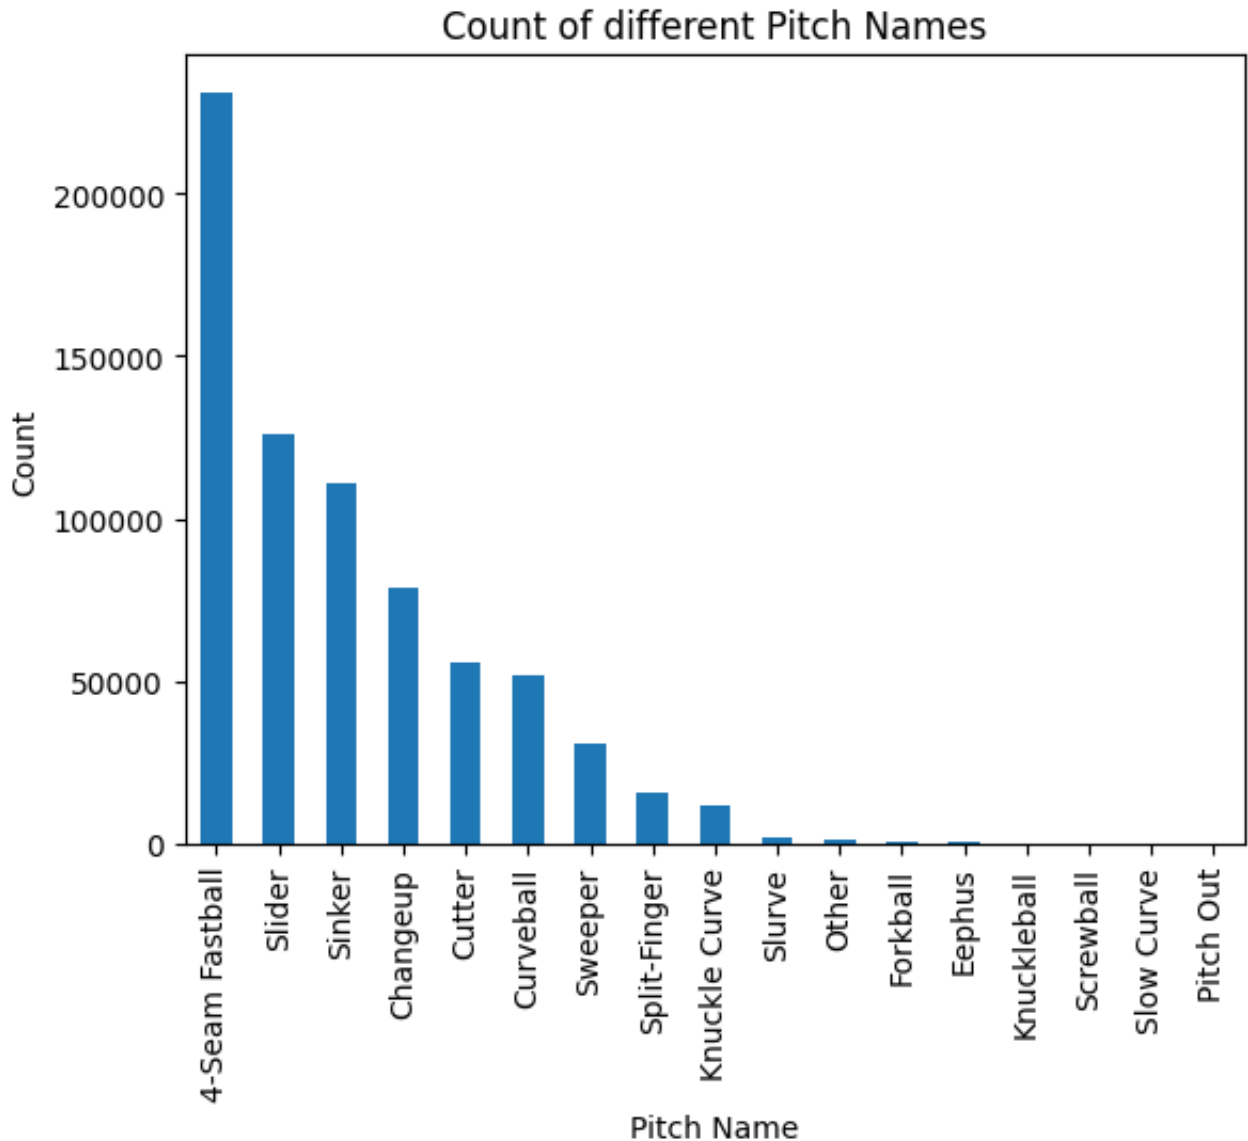
```
In [7]: #Deleting two rows that had pitch_type as NaN which might cause an issue later
```

```
df_reg_filtered = df_reg_filtered[~df_reg_filtered.pitch_name.isin([np.nan])]
```

```
In [8]: #First, we look at the different pitch types that these pitchers throw  
pitch_name_counts = df_reg['pitch_name'].value_counts()
```


```
In [9]: #We plot a count graph for the different types of pitches  
pitch_name_counts.plot(kind="bar")  
plot.title("Count of different Pitch Names")  
plot.xlabel("Pitch Name")  
plot.ylabel("Count")
```

```
Out[9]: Text(0, 0.5, 'Count')
```



```
In [10]: df_reg['type'].value_counts().plot(kind="bar")
```

```
Out[10]: <Axes: xlabel='type'>
```



```
In [11]: #Next, we divide the pitches for each individual pitcher

cross_tab_prop = pd.crosstab(index=df_reg_filtered['player_name'],
                              columns=df_reg_filtered['pitch_name'],
                              normalize="index")
cross_tab = pd.crosstab(index=df_reg_filtered['player_name'],
                        columns=df_reg_filtered['pitch_name'])
cross_tab
```

Out[11]:

pitch_name	4-Seam Fastball	Changeup	Curveball	Cutter	Knuckle Curve	Pitch Out	Sinker	Slider	(
player_name									
Eovaldi, Nathan	831	0	301	393	0	0	0	90	
Gallen, Zac	1601	453	0	333	738	0	12	111	
Gray, Sonny	766	188	473	370	0	0	442	0	
Javier, Cristian	1682	119	0	0	217	0	0	864	
Kelly, Merrill	842	648	308	508	0	0	342	155	
López, Pablo	1043	638	377	0	0	0	320	0	
Montgomery, Jordan	312	671	647	45	0	0	1245	0	
Nola, Aaron	904	390	0	230	977	0	586	0	
Pfaadt, Brandon	731	195	89	0	0	2	160	0	
Strider, Spencer	1826	225	0	0	0	1	0	1048	
Suárez, Ranger	434	372	384	253	0	0	567	13	
Valdez, Framber	0	462	718	385	0	0	1385	0	
Verlander, Justin	1302	121	520	0	0	0	0	659	
Wheeler, Zack	1370	13	361	453	0	0	571	0	

In [12]:

```
#Plotting the distribution of pitches for the 14 pitchers
cross_tab_prop.plot(kind='bar',
                    stacked=True,
                    colormap='tab10',
                    width=0.8,
                    figsize=(12, 6))

plot.title("Proportion of different Pitch types grouped by Player Name")
plot.xlabel("Player Name")
plot.ylabel("Proportion (Percentage and Count)")

for n, x in enumerate([*cross_tab.index.values]):
    for (proportion, count, y_loc) in zip(cross_tab_prop.loc[x],
```



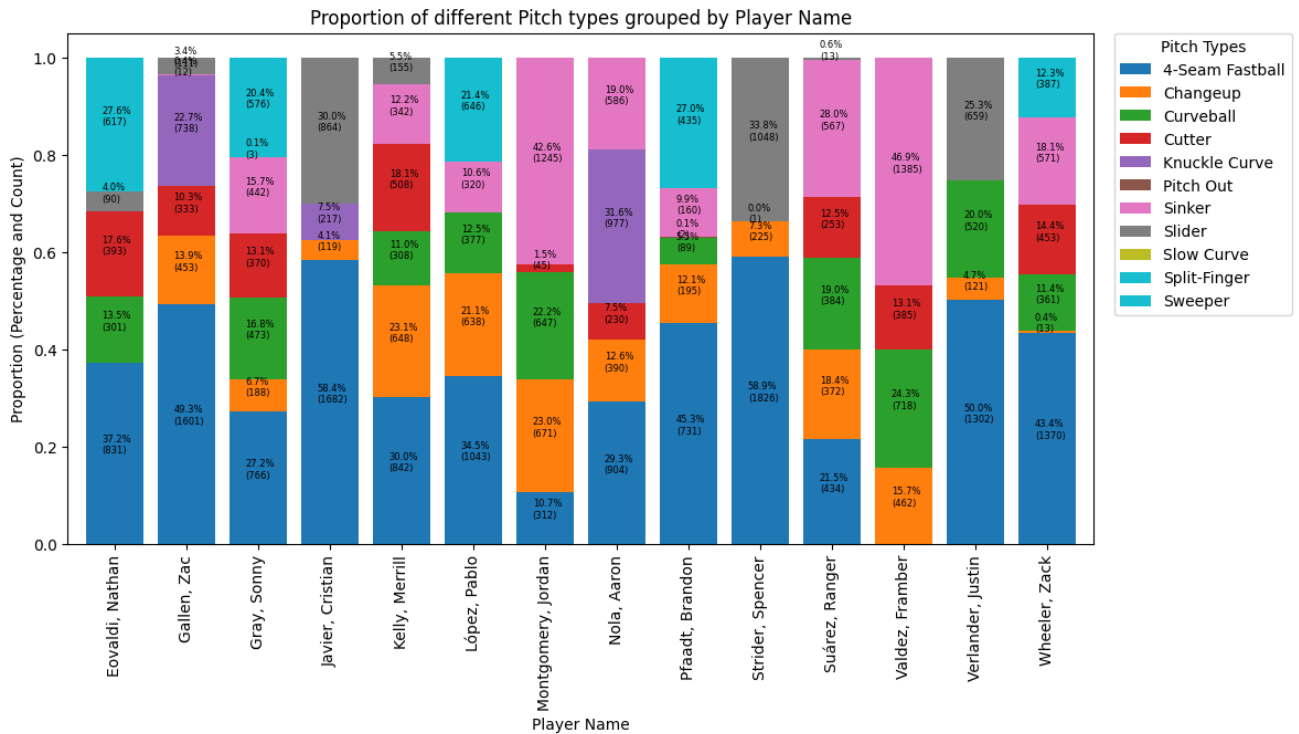
```

cross_tab.loc[x],
cross_tab_prop.loc[x].cumsum()):

if (proportion > 0):
    plot.text(x=n - 0.17,
              y=(y_loc - proportion) + (proportion / 2),
              s=f'{np.round(proportion * 100, 1)}%\n({count})',
              color="black",
              fontsize=6)

plot.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0, tit
plot.show()

```



```

In [13]: #We encode the data so that we can generate a heatmap for correlation of var
encoded_data = df_reg.copy()
le = LabelEncoder()
#Using Label Encoder, change all categorical data types to numerical
for col in ['pitch_name', 'player_name', 'stand', 'p_throws', 'home_team', '
            'type', 'pitch_name', 'if_fielding_alignment', 'of_fielding_al
            encoded_data[col] = le.fit_transform(df_reg[col])


```

```

In [14]: plot.figure(figsize=(8, 12))
heatmap = sns.heatmap(encoded_data.corr()[['pitch_name']].sort_values(by='pi
heatmap.set_title('Features Correlating with Pitch Name', fontdict={'fontsi


```

Features Correlating with Pitch Name



```
In [15]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 6))
sns.countplot(x='player_name', hue='pitch_name', data=df_post_filtered)
plt.title('Count of Different Pitch Types by Player Name (Strikeouts)')
plt.xlabel('Player Name')
plt.ylabel('Count')
```

```
plt.legend(title='Pitch Type', bbox_to_anchor=(1, 1))
plt.xticks(rotation=45, ha='right')
plt.show()
```



```
In [16]: #Divide the data into training and testing sets (regular season for the form
X_train, y_train = df_reg_filtered.iloc[:,df_reg_filtered.columns!= 'pitch_r
X_test, y_test = df_post_filtered.iloc[:,df_post_filtered.columns!= 'pitch_r
```

```
In [17]: print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(38456, 33)
(38456, 1)
(4532, 33)
(4532, 1)
```

```
In [18]: y_test
```

Out[18]:

	pitch_name
14	4-Seam Fastball
17	Slider
18	Slider
21	4-Seam Fastball
22	Knuckle Curve
...	...
987	Sinker
1048	4-Seam Fastball
1084	4-Seam Fastball
1119	4-Seam Fastball
1160	4-Seam Fastball

4532 rows × 1 columns

```
In [19]: #Using Label Encoder, change all categorical data types to numerical
from sklearn.preprocessing import MinMaxScaler
le = LabelEncoder()

for col in ['player_name', 'stand', 'p_throws', 'home_team', 'away_team',
            'type', 'if_fielding_alignment', 'of_fielding_alignment']:
    X_train.loc[:,col] = le.fit_transform(X_train[col])
    X_test.loc[:,col] = le.fit_transform(X_test[col])
X_train.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_train[['release_pos_x', 'release_pos_z']])
X_train.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_train[['release_pos_x', 'release_pos_z']])
X_test.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_test[['release_pos_x', 'release_pos_z']])
X_test.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_test[['release_pos_x', 'release_pos_z']])

y_train = le.fit_transform(np.ravel(y_train))
y_test = le.fit_transform(np.ravel(y_test))
```

```
In [20]: #Removing all nulls and filling them with 0. There are only 3 columns that contain nulls
X_train = X_train.fillna(0)
X_test = X_test.fillna(0)
```

```
In [21]: #Now, we create a covariance matrix that will be used for PCA
covar_matrix_train = np.matmul(X_train.T , X_train)

print ( "The shape of variance matrix = ", covar_matrix_train.shape)
```

The shape of variance matrix = (33, 33)

```
In [22]: #all columns are object types by default which causes an error in the eig function
covar_matrix_train = covar_matrix_train.apply(np.float64)
```

```
covar_matrix_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 33 entries, release_pos_x to of_fielding_alignment
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   release_pos_x                        33 non-null     float64
1   release_pos_z                        33 non-null     float64
2   player_name                         33 non-null     float64
3   batter                             33 non-null     float64
4   pitcher                             33 non-null     float64
5   stand                              33 non-null     float64
6   p_throws                            33 non-null     float64
7   home_team                           33 non-null     float64
8   away_team                           33 non-null     float64
9   type                                33 non-null     float64
10  balls                               33 non-null     float64
11  strikes                             33 non-null     float64
12  on_3b                               33 non-null     float64
13  on_2b                               33 non-null     float64
14  on_1b                               33 non-null     float64
15  outs_when_up                        33 non-null     float64
16  inning                             33 non-null     float64
17  game_pk                             33 non-null     float64
18  pitcher.1                           33 non-null     float64
19  fielder_2.1                         33 non-null     float64
20  fielder_3                           33 non-null     float64
21  fielder_4                           33 non-null     float64
22  fielder_5                           33 non-null     float64
23  fielder_6                           33 non-null     float64
24  fielder_7                           33 non-null     float64
25  fielder_8                           33 non-null     float64
26  fielder_9                           33 non-null     float64
27  at_bat_number                       33 non-null     float64
28  pitch_number                        33 non-null     float64
29  home_score                          33 non-null     float64
30  away_score                          33 non-null     float64
31  if_fielding_alignment               33 non-null     float64
32  of_fielding_alignment               33 non-null     float64
dtypes: float64(33)
memory usage: 9.8+ KB
```

```
In [23]: from scipy.linalg import eigh

values_train, vectors_train = eigh(covar_matrix_train, subset_by_index=(31,33))
print("Shape of eigen vectors = ",vectors_train.shape)
vectors_train = vectors_train.T
print("Updated shape of eigen vectors = ",vectors_train.shape)
```

```
Shape of eigen vectors = (33, 2)
Updated shape of eigen vectors = (2, 33)
```

```
In [24]: new_coordinates = np.matmul(vectors_train, X_train.T)
```

```
new_coordinates = np.vstack((new_coordinates, df_reg_filtered['pitch_name']))

# creating a new data frame for plotting the labeled points.
dataframe_pca = pd.DataFrame(data=new_coordinates, columns=["1st_principal",
print(dataframe_pca.head())
```

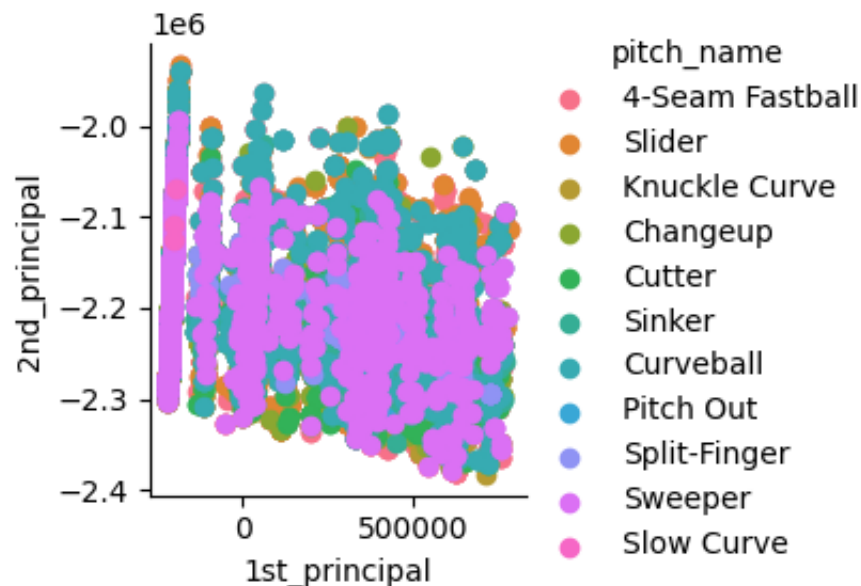
```
   1st_principal  2nd_principal  pitch_name
0  413464.045195 -2102242.806629  4-Seam Fastball
1  413464.045196 -2102242.806628  4-Seam Fastball
2  413464.045196 -2102242.806626  4-Seam Fastball
3  413464.045196 -2102242.806624  4-Seam Fastball
4  413464.045197 -2102242.806623  4-Seam Fastball
```

In [25]: #2-D representation of the PCA

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(16, 16))
sns.FacetGrid(dataframe_pca, hue="pitch_name").map(plt.scatter, '1st_principal', '2nd_principal')
plt.show()
```

<Figure size 1600x1600 with 0 Axes>



In [26]: #3-D representation of the PCA

```
from sklearn.decomposition import PCA

pca = PCA(n_components=11)
pca.fit(X_train)
X_pca = pca.transform(X_train)

ex_variance=np.var(X_pca,axis=0)
ex_variance_ratio = ex_variance/np.sum(ex_variance)
ex_variance_ratio
```

```

Xax = X_pca[:,0]
Yax = X_pca[:,1]
Zax = X_pca[:,2]


cdict = {0:'maroon',1:'olive', 2:'teal',3:'lightgreen',4:'black', 5:'red',6:
labl = {0:'Four-Seam Fastball',1:'Slider', 2:'Knuckle-Curve',3:'Changeup',4:

fig = plt.figure(figsize=(15,15))
ax = fig.add_subplot(111, projection='3d')

fig.patch.set_facecolor('white')
for l in np.unique(y_train):
    ix=np.where(y_train==l)
    ax.scatter(Xax[ix], Yax[ix], Zax[ix], c=cdict[l], s=40, label=labl[l])
# for loop ends
ax.set_xlabel("First Principal Component", fontsize=14)
ax.set_ylabel("Second Principal Component", fontsize=14)
ax.set_zlabel("Third Principal Component", fontsize=14)

ax.legend()
plt.show()

```



```
In [27]: from apyori import apriori
rules = apriori(transactions = np.array(df_reg_filtered.astype(str)), min_su
```

```
In [28]: def inspect(output):
    Left_Hand_Side = [tuple(result[2][0][0])[0] for result in output]
    support = [result[1] for result in output]
    confidence = [result[2][0][2] for result in output]
    lift = [result[2][0][3] for result in output]
    Right_Hand_Side = [tuple(result[2][0][1])[0] for result in output]
    return list(zip(Left_Hand_Side, support, confidence, lift, Right_Hand_Side))

output = list(rules)
output_data = pd.DataFrame(inspect(output), columns = ['Left_Hand_Side', 'Su
print(output_data)
```


	Left_Hand_Side	Support	Confidence	Lift	Right_Hand_Side
0	455117	0.107187	0.616236	5.560290	514888
1	455117	0.141746	0.814920	5.406930	547989
2	455117	0.171495	0.985947	5.399542	608324
3	455117	0.106797	0.613993	5.498770	643289
4	455117	0.161223	0.926895	5.456095	663656
...
5215	681082	0.113896	0.566404	4.088910	PHI
5216	681082	0.118525	0.589422	4.099622	PHI
5217	681082	0.189307	0.941420	4.126196	PHI
5218	681082	0.112986	0.561878	4.104001	PHI
5219	682998	0.180362	0.984388	4.443671	Standard

[5220 rows x 5 columns]

```
In [29]: # importing the required module
from mlxtend.preprocessing import TransactionEncoder
# initializing the transactionEncoder
te = TransactionEncoder()
te_ary = te.fit(np.array(df_reg_filtered.astype(str))).transform(np.array(df_reg_filtered))
dataset = pd.DataFrame(te_ary, columns=te.columns_)
# dataset after encoded
dataset
```

```
Out[29]:
```

	-0.41	-0.42	-0.43	-0.45	-0.46	-0.48	-0.49	-0.51	-0.53	-0.54	...	Swc
0	False	False	False	False	False	False	False	False	False	False	...	
1	False	False	False	False	False	False	False	False	False	False	...	
2	False	False	False	False	False	False	False	False	False	False	...	
3	False	False	False	False	False	False	False	False	False	False	...	
4	False	False	False	False	False	False	False	False	False	False	...	
...	
38451	False	False	False	False	False	False	False	False	False	False	...	
38452	False	False	False	False	False	False	False	False	False	False	...	
38453	False	False	False	False	False	False	False	False	False	False	...	
38454	False	False	False	False	False	False	False	False	False	False	...	
38455	False	False	False	False	False	False	False	False	False	False	...	

38456 rows x 1902 columns

```
In [30]: # importing the required module
from mlxtend.frequent_patterns import apriori, association_rules

# Extracting the most frequent itemsets via Mlxtend.
# The length column has been added to increase ease of filtering.
```

```
frequent_itemsets = apriori(dataset, min_support=0.1, use_colnames=True)
frequent_itemsets['length'] = frequent_itemsets['itemsets'].apply(lambda x:
# printing the frequent itemset
frequent_itemsets
```

Out[30]:

	support	itemsets	length
0	0.913226	(0)	1
1	0.914240	(1)	1
2	0.838153	(2)	1
3	0.509049	(3)	1
4	0.376742	(4)	1
...
15326	0.102299	(592663, 592206, 0, 607208, 681082, R, 664761,...	10
15327	0.103547	(665161, HOU, 1, 608324, <NA>, R, Standard, 66...	10
15328	0.101961	(592663, 592206, 607208, 1, 681082, R, Standar...	10
15329	0.103963	(665161, 663656, HOU, 1, 608324, R, Standard, ...	10
15330	0.102767	(592663, 592206, 607208, 1, 681082, R, 664761,...	10

15331 rows × 3 columns

```
In [31]: frequent_itemsets[(frequent_itemsets['length'] >= 2) &
(frequent_itemsets['support'] >= 0.2) & (frequent_itemsets['itemsets'].str.contains("4-Seam Fast
```

Out [31]:	support	itemsets	length
55	0.325671	(4-Seam Fastball, 0)	2
99	0.323955	(4-Seam Fastball, 1)	2
143	0.290514	(2, 4-Seam Fastball)	2
202	0.349282	(4-Seam Fastball, <NA>)	2
205	0.353157	(4-Seam Fastball, R)	2
...
8538	0.230419	(0, 1, 4-Seam Fastball, R, Standard, 2)	6
8784	0.285495	(0, 1, 4-Seam Fastball, R, Standard, <NA>)	6
9261	0.253276	(0, 4-Seam Fastball, R, Standard, <NA>, 2)	6
10195	0.249818	(1, 4-Seam Fastball, R, Standard, <NA>, 2)	6
12010	0.226493	(0, 1, 4-Seam Fastball, R, Standard, <NA>, 2)	7

63 rows × 3 columns

```
In [ ]: print(frequent_itemsets.nlargest(n = 5, columns = 'support'))
```


```
In [32]: #Trying out the elbow curve to get the number of important features/clusters

from sklearn.cluster import KMeans

wcss=[]

for i in range(1,11):
    kmeans = KMeans(n_clusters=i, init = 'k-means++', max_iter=300, n_init=10)
    kmeans.fit(X_train)
    wcss.append(kmeans.inertia_)

plt.plot(range(1,11),wcss)
plt.title('The Elbow Method Graph')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
```



MODELS

```
In [33]: from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score

model = MultinomialNB()

model.fit(X_train, y_train)

y_train_pred = model.predict(X_train)
y_train_prob = model.predict_proba(X_train)
print("accuracy:", accuracy_score(y_train, y_train_pred))
```

accuracy: 0.014042022051175369

```
In [34]: from sklearn.model_selection import GridSearchCV

param_grid = {'alpha': [0.0001, 0.001, 0.01, 0.1, 1.0]}
print("Parameter grid:\n{}".format(param_grid))

grid_search = GridSearchCV(MultinomialNB(), param_grid, cv=5,
                           return_train_score=True, scoring="accuracy")

grid_search.fit(X_train, y_train)
```

```
Parameter grid:
{'alpha': [0.0001, 0.001, 0.01, 0.1, 1.0]}
```

```
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/sklearn/model_selection/_split.py:737: UserWarning:
The least populated class in y has only 3 members, which is less than n_spli
ts=5.
```

```
warnings.warn(
```

Out [34]:

```
  ► GridSearchCV
    ► estimator: MultinomialNB
      ► MultinomialNB
```

```
In [35]: print("Best parameters: {}".format(grid_search.best_params_))
print("Best cross-validation score: {:.4f}".format(grid_search.best_score_))
```

```
Best parameters: {'alpha': 0.0001}
Best cross-validation score: 0.0141
```

```
In [36]: import multiprocessing
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, precision_score

nb = MultinomialNB(**grid_search.best_params_)
nb.fit(X_train, y_train)
test_score = nb.score(X_test, y_test)

y_pred = nb.predict(X_test)
y_pred_prob = nb.predict_proba(X_test)

print("Test accuracy:", accuracy_score(y_test, y_pred))
print("F1 score:", f1_score(y_test, y_pred, average='weighted'))
print("Precision score:", precision_score(y_test, y_pred, average='weighted'))
print("Recall score:", recall_score(y_test, y_pred, average='weighted'))
```

```
Test accuracy: 0.02581641659311562
F1 score: 0.010299713480654222
Precision score: 0.03380975153576563
Recall score: 0.02581641659311562
```

```
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/sklearn/metrics/_classification.py:1471: UndefinedMe
tricWarning: Precision is ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/sklearn/metrics/_classification.py:1471: UndefinedMe
tricWarning: Recall is ill-defined and being set to 0.0 in labels with no tr
ue samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [37]: from sklearn.ensemble import RandomForestClassifier

param_grid = {
```

```

    'bootstrap': [True, False],
    'max_depth': [11,15,17,19],
    'min_samples_leaf': [5, 10, 20],
    'min_samples_split': [10, 15, 20]
}
scorers = {
    'accuracy_score': make_scorer(accuracy_score)
}

rf_vec_gs = GridSearchCV(estimator=RandomForestClassifier(), param_grid=param_grid,
                        scoring=scorers, cv=10, n_jobs=multiprocessing.cpu_count(), verbose=0)

rf_vec_gs.fit(X_train, y_train)

rf_vec_best = rf_vec_gs.best_estimator_
rf_vec_best.fit(X_train, y_train)

print('Best model: %s' % str(rf_vec_gs.best_params_))
print('Best score: %f' % rf_vec_gs.best_score_)
print('Best test accuracy: %f' % rf_vec_best.score(X_test, y_test))

```

Fitting 10 folds for each of 72 candidates, totalling 720 fits

```

/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/sklearn/model_selection/_split.py:737: UserWarning:
The least populated class in y has only 3 members, which is less than n_splits=10.

```

```

    warnings.warn(
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/
python3.9/site-packages/joblib/externals/loky/process_executor.py:752: UserWarning:
A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

```

```

    warnings.warn(
Best model: {'bootstrap': True, 'max_depth': 15, 'min_samples_leaf': 20, 'min_samples_split': 15}
Best score: 0.486195
Best test accuracy: 0.442630

```

```

In [38]: from sklearn.neighbors import KNeighborsClassifier
        from sklearn.model_selection import GridSearchCV

        param_grid = {'n_neighbors': [2,4,7,11,15,18,20],
                       'metric': ["euclidean", "manhattan", "minkowski"]}

        knn_grid_search = GridSearchCV(estimator=KNeighborsClassifier(), param_grid=param_grid,
                                       return_train_score=True, verbose=1, scoring='accuracy')

        knn_grid_search.fit(X_train, y_train)
        print("Best parameters: {}".format(knn_grid_search.best_params_))
        print("Best cross-validation score: {:.2f}".format(knn_grid_search.best_score_))

```

Fitting 5 folds for each of 21 candidates, totalling 105 fits

```
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages/sklearn/model_selection/_split.py:737: UserWarning: The least populated class in y has only 3 members, which is less than n_splits=5.
```

```
warnings.warn(
```

```
Best parameters: {'metric': 'manhattan', 'n_neighbors': 20}
```

```
Best cross-validation score: 0.30
```

```
In [39]: optimal_knn_model = KNeighborsClassifier(**knn_grid_search.best_params_)
         optimal_knn_model.fit(X_train, y_train)

         y_pred = optimal_knn_model.predict(X_test)
```

```
In [40]: from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, precision_score, recall_score

         print("Test accuracy:", accuracy_score(y_test, y_pred))
         print("F1 score:", f1_score(y_test, y_pred, average='weighted'))
         print("Precision score:", precision_score(y_test, y_pred, average='weighted'))
         print("Recall score:", recall_score(y_test, y_pred, average='weighted'))
```

```
Test accuracy: 0.3287731685789938
```

```
F1 score: 0.2701272552247107
```

```
Precision score: 0.26569524377909404
```

```
Recall score: 0.3287731685789938
```

```
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/Users/kishan/Documents/Documents/IST 707 Applied Machine Learning/venv/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [42]: import optuna
         from sklearn.svm import SVC

         def objective(trial):
             # Define the hyperparameters to search
             C = trial.suggest_float('C', 1e-4, 1e4)
             gamma = trial.suggest_float('gamma', 1e-4, 1e4)
             kernel = trial.suggest_categorical('kernel', ['rbf'])

             # Create the SVM classifier with the suggested hyperparameters
             clf = SVC(C=C, gamma=gamma, kernel=kernel)

             # Train the classifier and evaluate on the validation set
             clf.fit(X_train, y_train)
             accuracy = clf.score(X_test, y_test)

             return accuracy
```

```
# Run the optimization
study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=10)

# Print the best hyperparameters and the corresponding accuracy
best_params = study.best_params
best_accuracy = study.best_value
print('Best Hyperparameters:', best_params)
print('Best Accuracy:', best_accuracy)
```

```
[I 2023-12-03 22:27:26,837] A new study created in memory with name: no-name-
-a0d3ebd5-153c-4932-9d81-775dd0484a67
[I 2023-12-03 22:32:29,282] Trial 0 finished with value: 0.35017652250661957
and parameters: {'C': 9178.370146770083, 'gamma': 4423.787399043729, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 22:37:36,796] Trial 1 finished with value: 0.35017652250661957
and parameters: {'C': 8887.906518984144, 'gamma': 9781.59198639811, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 22:42:39,127] Trial 2 finished with value: 0.35017652250661957
and parameters: {'C': 3738.4275754521464, 'gamma': 8313.346516779042, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 22:47:41,349] Trial 3 finished with value: 0.35017652250661957
and parameters: {'C': 1624.1860998197994, 'gamma': 4212.212615812341, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 22:52:43,277] Trial 4 finished with value: 0.35017652250661957
and parameters: {'C': 618.391921643613, 'gamma': 8542.15187245932, 'kernel':
'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 22:57:43,648] Trial 5 finished with value: 0.35017652250661957
and parameters: {'C': 6730.700106298582, 'gamma': 4222.483566343328, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 23:02:44,250] Trial 6 finished with value: 0.35017652250661957
and parameters: {'C': 1411.6483741007348, 'gamma': 6280.727288509174, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 23:07:44,788] Trial 7 finished with value: 0.35017652250661957
and parameters: {'C': 5776.239048814928, 'gamma': 9505.940407374144, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 23:12:45,573] Trial 8 finished with value: 0.35017652250661957
and parameters: {'C': 6005.036674308265, 'gamma': 3566.0572289026345, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
[I 2023-12-03 23:17:45,960] Trial 9 finished with value: 0.35017652250661957
and parameters: {'C': 9931.256496019494, 'gamma': 2799.1959978800214, 'kerne
l': 'rbf'}. Best is trial 0 with value: 0.35017652250661957.
Best Hyperparameters: {'C': 9178.370146770083, 'gamma': 4423.787399043729, '
kernel': 'rbf'}
Best Accuracy: 0.35017652250661957
```

INDIVIDUAL PITCHERS

```
In [126... nola_reg_filtered = df_reg_filtered[df_reg_filtered['pitcher'].isin([605400]
nola_post_filtered = df_post_filtered[df_post_filtered['pitcher'].isin([6054

gallen_reg_filtered = df_reg_filtered[df_reg_filtered['pitcher'].isin([66867
gallen_post_filtered = df_post_filtered[df_post_filtered['pitcher'].isin([66
```



```
montgomery_reg_filtered = df_reg_filtered[df_reg_filtered['pitcher'].isin([6
montgomery_post_filtered = df_post_filtered[df_post_filtered['pitcher'].isin
```

```
In [127... X_train_nola, y_train_nola = nola_reg_filtered.iloc[:,nola_reg_filtered.colu
X_test_nola, y_test_nola = nola_post_filtered.iloc[:,nola_post_filtered.colu

X_train_gallen, y_train_gallen = gallen_reg_filtered.iloc[:,gallen_reg_filte
X_test_gallen, y_test_gallen = gallen_post_filtered.iloc[:,gallen_post_filte

X_train_montgomery, y_train_montgomery = montgomery_reg_filtered.iloc[:,mont
X_test_montgomery, y_test_montgomery = montgomery_post_filtered.iloc[:,montg
```

```
In [128... #Using Label Encoder, change all categorical data types to numerical
from sklearn.preprocessing import MinMaxScaler
le = LabelEncoder()

for col in ['player_name', 'stand', 'p_throws', 'home_team', 'away_team',
            'type', 'if_fielding_alignment', 'of_fielding_alignment']:
    X_train_nola.loc[:,col] = le.fit_transform(X_train_nola[col])
    X_test_nola.loc[:,col] = le.fit_transform(X_test_nola[col])

    X_train_gallen.loc[:,col] = le.fit_transform(X_train_gallen[col])
    X_test_gallen.loc[:,col] = le.fit_transform(X_test_gallen[col])

    X_train_montgomery.loc[:,col] = le.fit_transform(X_train_montgomery[col])
    X_test_montgomery.loc[:,col] = le.fit_transform(X_test_montgomery[col])

X_train_nola.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_train_r
X_train_nola.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_train_r
X_test_nola.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_test_nola
X_test_nola.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_test_nola
X_train_nola = X_train_nola.fillna(0)
X_test_nola = X_test_nola.fillna(0)
y_train_nola = le.fit_transform(np.ravel(y_train_nola))

X_train_gallen.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_train
X_train_gallen.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_train
X_test_gallen.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_test_g
X_test_gallen.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_test_g
X_train_gallen = X_train_gallen.fillna(0)
X_test_gallen = X_test_gallen.fillna(0)
y_train_gallen = le.fit_transform(np.ravel(y_train_gallen))

X_train_montgomery.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_t
X_train_montgomery.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_t
X_test_montgomery.loc[:, 'release_pos_x'] = MinMaxScaler().fit_transform(X_te
X_test_montgomery.loc[:, 'release_pos_z'] = MinMaxScaler().fit_transform(X_te
X_train_montgomery = X_train_montgomery.fillna(0)
X_test_montgomery = X_test_montgomery.fillna(0)
y_train_montgomery = le.fit_transform(np.ravel(y_train_montgomery))
```

```
In [143... X_test_nola.iloc[2]
```

```
Out[143... release_pos_x      0.724638
release_pos_z      0.673077
player_name         0.0
batter              682998.0
pitcher             605400.0
stand               0.0
p_throws            0.0
home_team           0.0
away_team           1.0
type                2.0
balls               2.0
strikes             2.0
on_3b               0.0
on_2b               0.0
on_1b               0.0
outs_when_up        1.0
inning              5.0
game_pk             748539.0
pitcher.1           605400.0
fielder_2.1         592663.0
fielder_3           547180.0
fielder_4           681082.0
fielder_5           664761.0
fielder_6           607208.0
fielder_7           669016.0
fielder_8           679032.0
fielder_9           592206.0
at_bat_number       37.0
pitch_number        8.0
home_score          1.0
away_score          3.0
if_fielding_alignment 1.0
of_fielding_alignment 0.0
Name: 12, dtype: Float64
```

```
In [139... #Trying to predict a pitch type based on some input
pitch_to_predict = X_test_nola.iloc[2]

rf_nola = rf_vec_best.fit(X_train_nola.values, y_train_nola)
print("Nola: %s" % (df_reg_filtered['pitch_name'].unique()[rf_nola.predict([

rf_gallen = rf_vec_best.fit(X_train_gallen.values, y_train_gallen)
print("Gallen: %s" % (df_reg_filtered['pitch_name'].unique()[rf_gallen.predi

rf_montgomery = rf_vec_best.fit(X_train_montgomery.values, y_train_montgomer
print("Montgomery: %s" % (df_reg_filtered['pitch_name'].unique()[rf_montgome
```

Nola: Changeup
Gallen: 4-Seam Fastball
Montgomery: Cutter